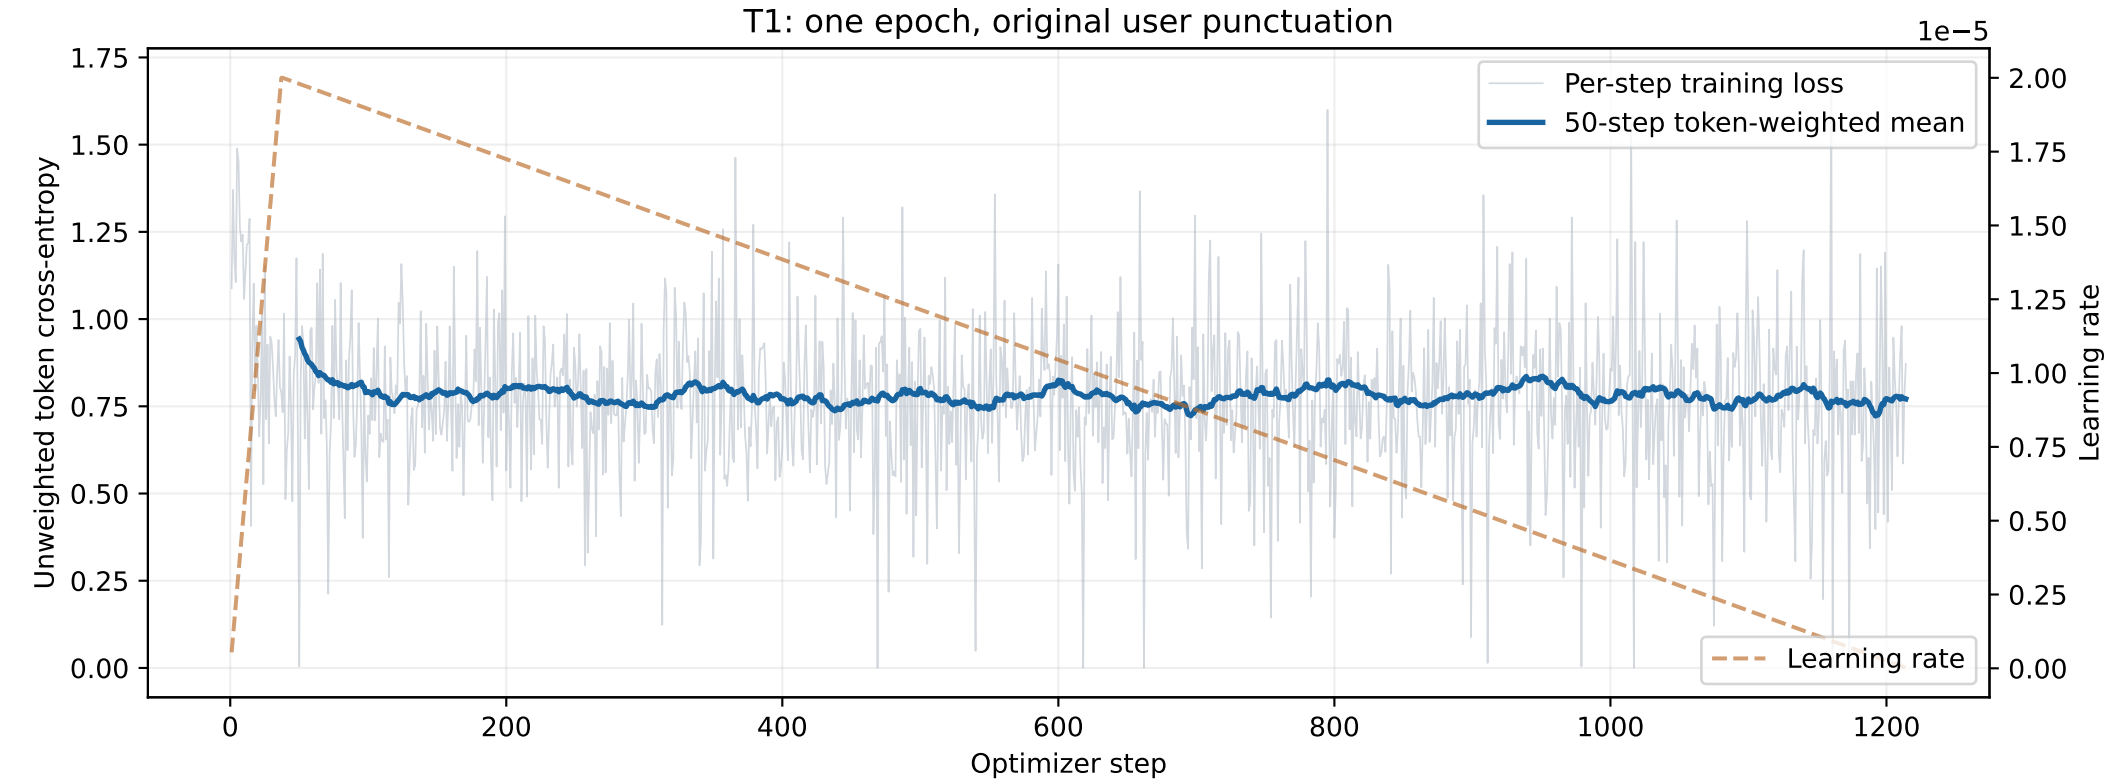
T1: one epoch, original user punctuation
1e−5
Per-step training loss
50-step token-weighted mean
Learning rate
Unweighted token cross-entropy
Learning rate
Optimizer step
1.75
1.50
1.25
1.00
0.75
0.50
0.25
0.00
2.00
1.75
1.50
1.25
1.00
0.75
0.50
0.25
0.00
0
200
400
600
800
1000
1200

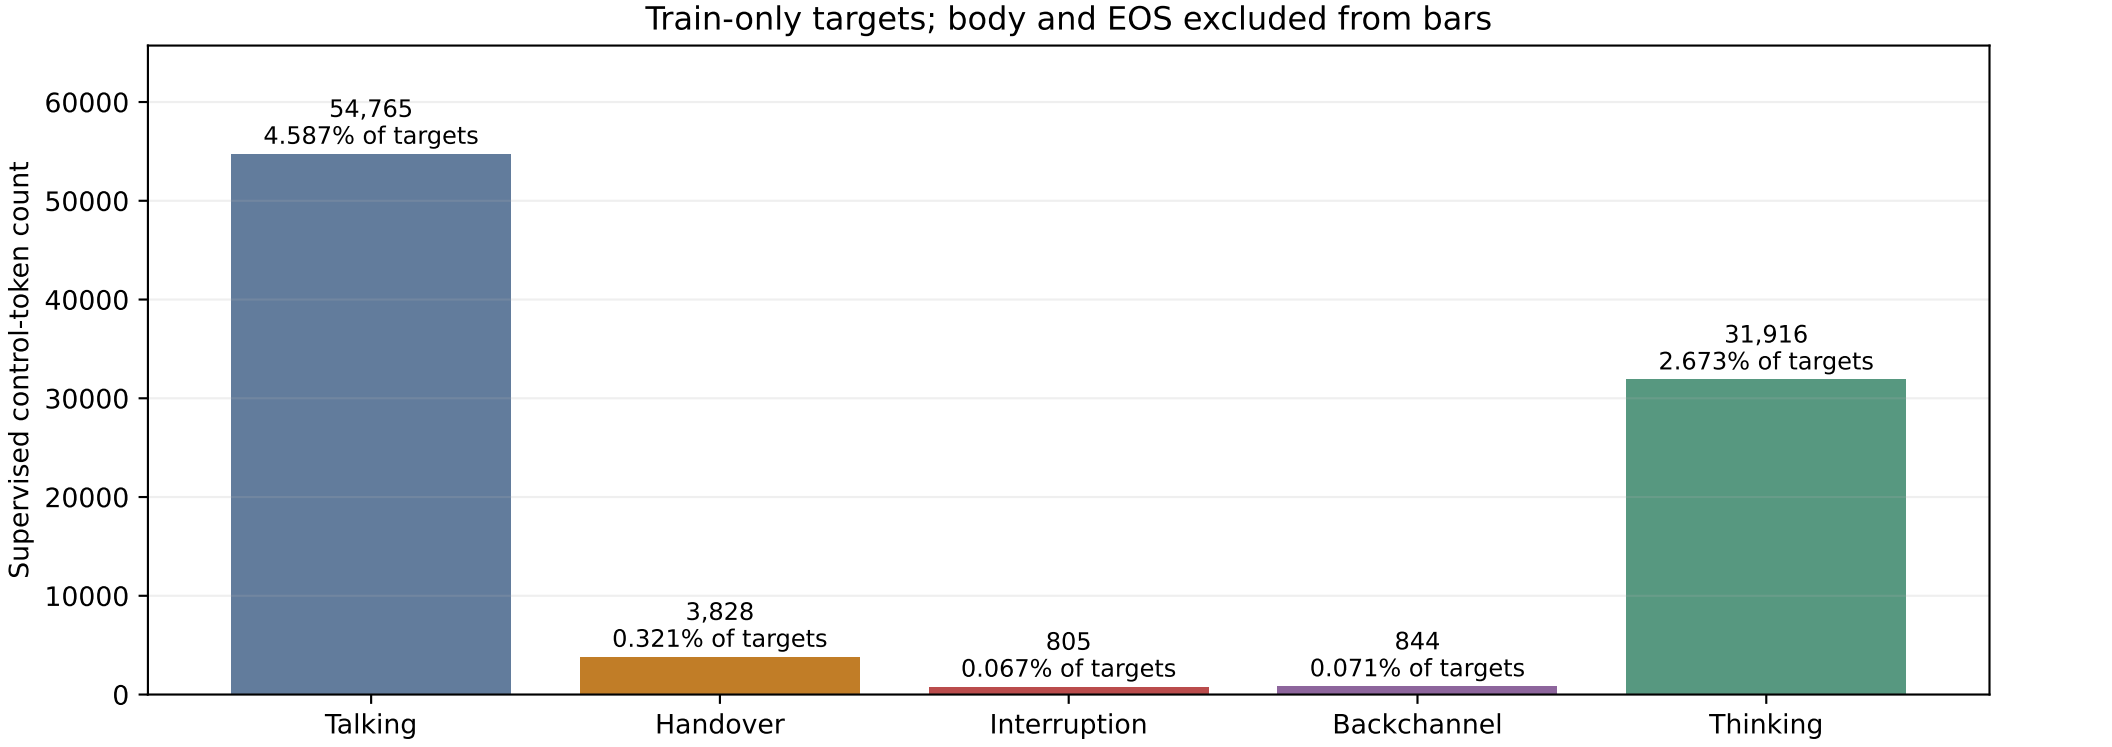
Train-only targets; body and EOS excluded from bars
Supervised control-token count
60000
50000
40000
30000
20000
10000
0
54,765
4.587% of targets
3,828
0.321% of targets
805
0.067% of targets
844
0.071% of targets
31,916
2.673% of targets
Talking
Handover
Interruption
Backchannel
Thinking